CS 31: Algorithms (Spring 2019): Lecture 11

Date: 30th April, 2019
Topic: Graph Algorithms 1: Depth First Search
Disclaimer: These notes have not gone through scrutiny and in all probability contain errors.
Please notify errors on Piazza/by email to deeparnab@dartmouth.edu.

---

1   Depth First Search (DFS)

We start graph algorithms with the pretty intuitive, but surprisingly powerful, *depth first search* (DFS). This algorithm solves the reachabilty problem, but then in one swoop solves much more. It also runs in $\Theta(n + m)$ time. Let's get to it.

1.1   DFS from a vertex

The objective for this section is to solve the reachability problem: given vertices u and v, is there a path from u to v in G? The algorithm is the same for solving mazes: we start at vertex u first planting a flag there; subsequently, we visit any unflagged neighbor x of u unraveling a thread to remembering we came to x from u, and repeat the procedure from x; if at any time we are stuck with flags all around, then we ravel the thread back to the place where we came from. Formally, we use recursion.

```
 1: global visited[1 : n] initialized to 0.
 2: global Tree F initialized to ⊥.
 3: procedure DFS(G, v):    ▷ We assume V = {1, 2, ..., n}
 4:        ▷ Returns a tree F rooted at v
 5:        visited[v] ← 1.  ▷ Mark v visited.
 6:        Add vertex v to F.
 7:        for u neighbor of v do:   ▷ In an arbitrary, but fixed order
 8:            if visited[u] = 0 then:
 9:                Add edge (v, u) to the tree F.
10:                DFS(G, u)
11:        return F
```

Exercise: *Please draw a directed graph G and run the above pseudocode by hand to get a feel for it.*

Once we run DFS(G, v), we obtain the object F. The next claim tells us that F contains all the vertices reachable from v.

Claim 1. Upon running DFS(G, v), we have visited$[x] = 1$, or equivalently $x \in F$, if and only if x is reachable from v. Furthermore, x is a descendant of v in F in this case.

Although the above may seem obvious, it does require a formal proof. And we provide one below.

Proof. First we prove if visited$[x] = 1$, then x is reachable from v, in fact, using edges in F. The proof is by induction on the *time* at which visited$[x]$ was set to 1. Imagine every time the algorithm runs Line 5, we increment time by 1. At time 0, this was set visited$[v] = 1$ and v is reachable from v in F. Now pick a vertex x whose visited$[x] = 1$ is set at time t. This happens because of some $y \in V$ such that (a) $(y, x) \in E$, and (b) the run of DFS(G, y) calls DFS(G, x). In that case, (a) visited$[y] = 1$ has been set strictly before time t implying, by Induction, y is reachable from v in F, and (b) we add edge (y, x) to F, which then implies x is reachable from v in F.

Now for the other direction. Suppose there exists a vertex x which is reachable from v in G but visited$[x] = 0$. Since x is reachable from v, there is a path $(v = v_0, v_1, \ldots, v_k = x)$ in G. Let us pick the *last* vertex v_i in this path which has visited$[v_i] = 1$; clearly $0 \le i < k$. Since visited$[v_i] = 1$, we have run DFS(G, v_i). But the for-loop in the algorithm would then call DFS(G, v_{i+1}) since visited$[v_{i+1}] = 0$. But that would set visited$[v_{i+1}] = 1$, and once visited a vertex is never "un-visited". This is a contradiction, and thus all vertices x reachable from v have visited$[x] = 1$. $\square$

Claim 2. F is a tree.

✎

> **Exercise:** *Prove the above formally like the previous claim. Practice writing proofs.*

Theorem 1. The REACHABILITY problem can be solved in $\Theta(n + m)$ time.

Proof. To solve the reachability problem, we just run DFS from u and check if visited$[v] = 1$ or not. To get the path, we can use the tree F. The running time is $\Theta(n + m)$ since every edge is considered in the for-loop at most twice (once if G is directed). To see it in a "different" way, the running time in the for-loop at most the sum of degrees (out-degrees, if directed). ▢

1.2 DFS on the whole graph

The next algorithm is a *traversal* over all vertices of the graphs using the subroutine DFS(G, v). This is called the *depth first traversal* algorithm of the graph G, but is also called the depth first search of G. The input to this algorithm is the graph G and a permutation σ of the vertices. This permutation tells the algorithm the order in which to "explore" vertices, that is, to run DFS(G, v).

The output to this algorithm has a lot of things; these objects contain surprising amounts of information about G, as we will see below.

- One output is a couple of vectors first$[1:n]$ and last$[1:n]$ where for any vertex v, first$[v]$ notes the "time" at which the algorithm starts exploring from v, that is, DFS(G,v) is called, and last$[v]$ denotes the "time" the exploring ends, that is, the subroutine DFS(G,v) terminates.

- The other output is a *forest F* spanning all the vertices of G. Each tree in the forest is rooted, and is assumed to be directed (even when G is not) away from the root. Together with this we store the scalar fcomp which counts the number of trees in F, the array root$[1:$ fcomp$]$ where root(i) will store the root of the ith tree in F, and the array Fcomp$[1:n]$ where Fcomp$[v]$ contains a number between 1 and fcomp indicating the tree in which v exists.

The algorithm is simple: it has a for-loop going over all vertices in the order σ; if the vertex is unvisited, then we run DFS(G,v) on it starting a new tree rooted from v. We end when there are no more vertices left. Here is the full pseudocode, where we have enhanced DFS(G,v) to take care of what we need.

```
1:  procedure DFS(G, σ[1 : n]): ▷ σ is an ordering of the vertices
2:      global array visited[1 : n] initialized to all 0.
3:      global array first[1 : n], last[1 : n], root[1 : n], Fcomp[1 : n] initialized to all 0.
4:      global scalar fcomp, time initialized to 0.
5:      global Forest F initialized to ∅.

6:      for v in σ do:
7:          if visited[v] = 0 then:▷ v hasn't been visited yet:
8:              fcomp ← fcomp + 1 ▷ Increase the number of trees in the forest
9:              root[fcomp] ← v ▷ Set v to be the root of the new tree
10:             DFS(G, v)

11: procedure DFS(G, v):
12:     visited[v] ← 1;  Add v to F.
13:     Fcomp[v] ← fcomp. ▷ Set v's tree in the forest
14:     time ← time + 1.
15:     Set first[v] ← time. ▷ Start exploring.
16:     for u neighbor of v do: ▷ In an arbitrary, but fixed order
17:         if visited[u] = 0 then:
18:             Add edge (v, u) to the forest F.
19:             ▷ It will be added to the fcompth component.
20:             DFS(G, u)
21:     time ← time + 1.
22:     Set last[v] ← time.
```

Claim 3. The running time of DFS(G,σ) is $\Theta(m+n)$ for any σ.

3

Figure 1: The edges that appear in the forest are marked in solid, while the remaining edges are dotted. The first and last are noted near the vertices. In the third figure on the right, the interval is the $[\text{first}[v], \text{last}[v]]$ interval.
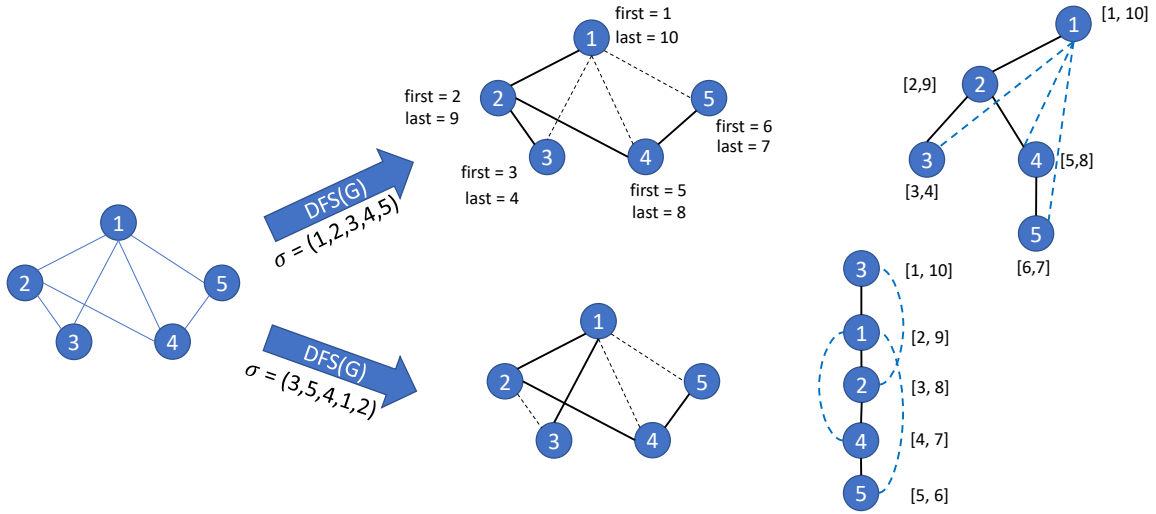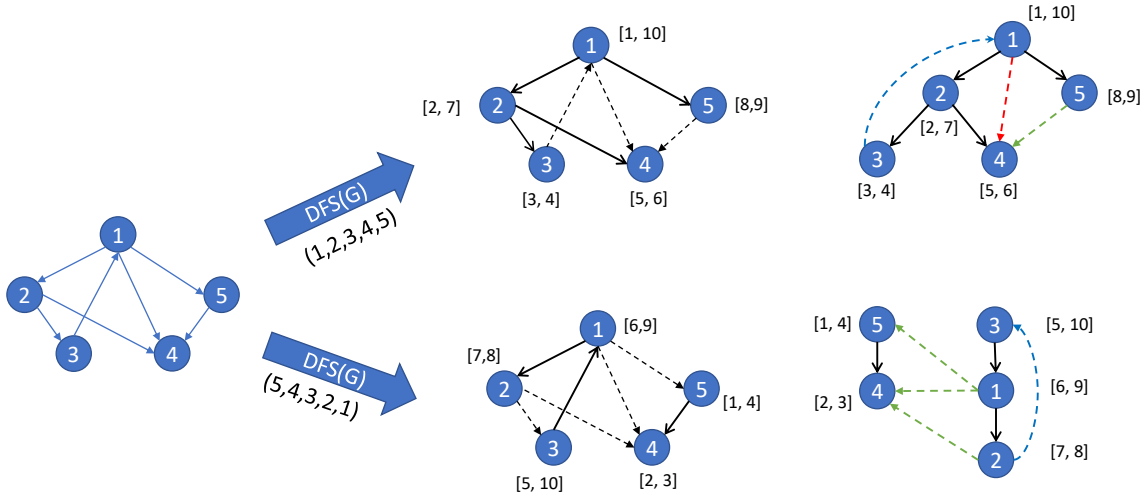

Figure 2: The edges that appear in the forest are marked in solid, while the remaining edges are dotted. The first and last are noted near the vertices. In the third figure on the right, the interval is the $[\text{first}[v], \text{last}[v]]$ interval.

F / DFS Trees and forests are not unique.
Different permutations can lead to different outcomes: see Figure 1 and Figure 2.

Edge Classification. Running DFS on a graph G with any ordering σ leads to four kinds of edges.

Edges that go directly from a parent to a child.

- (Forest Edges.) These are the edges present in F. These are marked black and solid in the Figures.

Leads to previously visited vertex parent -> child, gramdparent -> child etc.

- (Back Edges.) These edges go from a descendant to an ancestor. These are marked blue and dotted.
- (Forward Edges.) These edges go from an ancestor to a descendant. These are marked red and dotted. For undirected graphs the forward edges are all back edges (there is no direction).

Edges crossing to siblings or other trees in the forest

- (Cross Edges.) All the rest. They can be among pairs in the same component, or not. These are marked green and dotted.

Properties. Next, we state and prove **three** properties of the output of the DFS algorithm. Before reading the proofs, it will be useful to check if the statement is indeed correct by looking at the examples in Figures 1 and 2 (or any other figures you have privately made). For all the properties below, we assume we have run $\text{DFS}(G, \sigma)$ for some arbitrary ordering σ.

Lemma 1 (Nested Interval Property). For any two vertices u and v, with $\text{first}[u] < \text{first}[v]$, exactly one of the following two properties hold.

DFS started exploring u before v and terminated on v before terminating on u

- $\text{first}[u] < \text{first}[v] < \text{last}[v] < \text{last}[u]$ *and* v is a descendant of u in F.
- $\text{first}[u] < \text{last}[u] < \text{first}[v] < \text{last}[v]$ *and* neither is a descendant of the other.

DFS explored and terminated on u before beginning to explore v

This shows that the n intervals of the form $[\text{first}[v], \text{last}[v]]$ don't "criss-cross" (although one may be contained in the other). This property is called the *nested* property (also called laminar property).

Proof. Since $\text{first}[u] < \text{first}[v]$, we call $\text{DFS}(G, u)$ before we call $\text{DFS}(G, v)$. If v is ever discovered in the run of $\text{DFS}(G, u)$, then (a) it will be a descendant of u in F (by Claim 1), and (b) we must have $\text{last}[v] < \text{last}[u]$ since this recursive call must end before u's recursive call ends. This is case 1.

The only other case is v has not been discovered in the run of $\text{DFS}(G, u)$. In that case, $\text{last}[u] < \text{first}[v]$ by definition. Furthermore, v is not a descendant of u, and u can't be a descendant of v since v hasn't been even discovered yet. This is case 2. $\square$

The above property is useful, and will be useful in proving some other properties below. But it also allows us to classify the edges (not in the forest F) just looking at the first and the last values.

- (Back Edges.) Edges $(u, v) \in E \setminus F$ with $\text{first}[v] < \text{first}[u] < \text{last}[u] < \text{last}[v]$.

5

$U \rightarrow V$

- (Forward Edges.) Edges $(u, v) \in E \setminus F$ with $\text{first}[u] < \text{first}[v] < \text{last}[v] < \text{last}[u]$.
- (Cross Edges.) Edges $(u, v) \in E$ such that the intervals $[\text{first}[u], \text{last}[u]]$ and $[\text{first}[v], \text{last}[v]]$ are disjoint.

Lemma 2 (Edge Property). Let (u, v) be any edge in G with $\text{first}[u] < \text{first}[v]$. Then, we must have $\text{last}[v] < \text{last}[u]$.

Proof. Suppose not. By the Nested Interval Property, we must have $\text{first}[u] < \text{last}[u] < \text{first}[v] < \text{last}[v]$. That happens when $\text{DFS}(G, u)$ terminates before visited$[v]$ is set to 1. But the Line 16 would discover v contradicting the above. □

We are now ready for the first application of DFS – we can solve CONNECTED COMPONENTS of an Undirected Graph using the following lemma.

Lemma 3. Let $G = (V, E)$ be any undirected graph and consider the forest F returned by $\text{DFS}(G, \sigma)$ with any permutation σ. The components of F are precisely the connected components of G.

Proof. Let $V_1, \ldots, V_k$ be the vertices in the various trees of the forest F. Clearly $G[V_i]$ is connected since they are connected in the forest. We claim that there is no edge of the form (u, v) with $u \in V_i$ and $v \in V_j$. Suppose there is, and without loss of generality assume $\text{first}[u] < \text{first}[v]$ (this is where we are using the undirectedness of G). By the edge property, we have $\text{first}[u] < \text{first}[v] < \text{last}[v] < \text{last}[u]$. But this means v is a descendant of u in F contradicting the fact they exist in different connected components of F. □

Theorem 2. CONNECTED COMPONENTS of an undirected graph can be found in $\Theta(n+m)$ time by running $\text{DFS}(G, \sigma)$ for any ordering σ.

Moving on to more properties.

Lemma 4 (Path Property). If $(u = v_1, v_2, \ldots, v_k = v)$ is a path in G from u to v such that $\text{first}[u] < \text{first}[v_i]$ for all $2 \leq i \leq k$, then, $\text{last}[v_i] < \text{last}[u]$ for all $2 \leq i \leq k$.

In English, if there is a path from a vertex u to a vertex v such that u is the first vertex to be discovered among them, then all the vertices in the path are descendants of u in the DFS forest.

Proof. Suppose not. Choose the smallest $2 \leq i \leq k$ for which $\text{last}[u] < \text{last}[v_i]$. By the choice of i, we get $\text{last}[v_{i-1}] < \text{last}[u]$. Also note (v_{i-1}, v_i) is an edge.
　Case 1: $\text{first}[v_{i-1}] < \text{first}[v_i]$. In this case, the Edge Property would imply $\text{last}[v_i] < \text{last}[v_{i-1}]$, and thus $\text{last}[v_i] < \text{last}[u]$. Which is what we supposed wasn't true.
　Case 2: $\text{first}[v_i] < \text{first}[v_{i-1}]$. In that case, we see $\text{first}[u] < \text{first}[v_i] < \text{last}[u] < \text{last}[v_i]$ which violates the Nested Interval Property. □

The above property allow us immediately to solve the CYCLE? problem. The following theorem implies the algorithm: run $\text{DFS}(G, \sigma)$ and check if any of the edges is a *back edge* (which is one linear time scan over all the edges and checking the first and the last).

6

Lemma 5. A graph G is **acyclic** if and only if there are no back edges.

Proof. One direction is trivial – if G has a back edge, then there is clearly a cycle. If the back-edge is (u, v), then by definition there is a path from v to u using F-edges, and then take the (u, v) edge back.

The other direction is more interesting. If G has a cycle C with k vertices $(v_1, \ldots, v_k, v_1)$, then without loss of generality let v_1 be the vertex with the smallest first$[v_i]$ in this cycle. Since there is a path from v_1 to v_k, using the Path property and the fact that first$[v_1]$ is the smallest, we get first$[v_1] <$ first$[v_k] <$ last$[v_k] <$ last$[v_1]$. But this implies (v_k, v_1) is a back-edge. $\square$

Theorem 3. CYCLE? can be solved in $\Theta(n + m)$ time using DFS.